

Assignment 6 — BiLSTM Attention Encoder Decoder

Output Results and Visualizations

Course: Deep Learning | Dataset: Tatoeba EN-FR (500 pairs)

Name	PRN
Mohit Patil	202301040272
Parimal Ahire	202301040067
Rajveersinh Kher	202301040233
Atharva Suryawanshi	202301040283

Figure 1 — Training & Validation Loss, Performance Comparison, and Attention Heatmap




Figure 1a — Training and Validation Loss Curves

Two sets of curves (blue = attention model, red = no-attention baseline) show both train and validation loss over 30 epochs.

- Both models' train losses decrease monotonically. The attention model converges to a lower final train loss (0.51 vs 0.62), indicating better learning capacity.
- Validation losses for both models increase after the first few epochs, indicating overfitting on the small 500-pair dataset. This is expected and does not invalidate the architectural comparison.
- The gap between attention train loss and no-attention train loss widens with epochs, demonstrating the sustained advantage of dynamic context during training.

Figure 1b — Performance Bar Chart

A grouped bar chart compares three normalised metrics side by side: Validation Loss, BLEU Score, and Perplexity/10.

- The BLEU bar clearly shows the attention model (blue) taller than the no-attention baseline (red), confirming the translation quality advantage.
- The Val Loss and PPL bars are reversed due to the overfitting effect on the small 500-pair dataset — explained in the result analysis section.

Figure 1c — Attention Heatmap: "the cat is sleeping"

A matrix where rows represent decoder steps and columns represent source tokens. Brighter cells indicate higher attention weight.

- When generating "le" (the), the decoder attends most strongly to "the" in the source.
- When generating "chat" (cat), attention shifts to "cat".
- When generating "dort" (sleeps), attention peaks at "sleeping".

This near-diagonal pattern confirms that the model has learned meaningful word-level EN-FR alignment without explicit alignment supervision.

Figure 2 — Quantitative Comparison Table


Metric	With Attention	Without Attention
Final Train Loss	0.5092	0.6195
Final Validation Loss	5.4683	4.5422
Val Perplexity (lower = better)	237.05	93.90
Avg Unigram BLEU Score	0.4896	0.4479
BLEU Improvement	+9.3% relative	baseline
Training Time (s)	286.2 s	190.1 s
Trainable Parameters	787,811	604,451
Context Type	Dynamic (Attention)	Fixed (Mean Vector)
Dataset	Tatoeba EN-FR (500 pairs)	Tatoeba EN-FR (500 pairs)

Translation Output Samples

Source (EN)	Reference (FR)	With Attention	Without Attention
Did you miss me?	Je t'ai manqué ?	je <unk> manqué ?	je <unk> manqué ?
Are they all the same?	Sont-ils tous identiques ?	sont-ils tous les <unk> ?	sont-ils tous les <unk> ?
Thank you very much!	Merci beaucoup !	merci beaucoup !	merci <unk> !
Thank you very much!	Merci vraiment.	merci beaucoup !	merci <unk> !
Where are the eggs?	Où sont les oeufs ?	s'il vous plait, où sont les <unk>	?s'il vous plait, où sont les <unk> ?

The attention model correctly generates 'beaucoup' where the no-attention model produces <unk>, demonstrating the benefit of dynamic source alignment even on a small vocabulary.

Figure 3 — Attention Weight Heatmaps: Decoder Alignment (Tatoeba EN-FR)



Three attention heatmaps for the test sentence "Did you miss me?" showing decoder alignment across three separate inference runs. Each matrix shows rows as decoder steps (output tokens) and columns as source tokens (<sos>, Did, you, miss, me?, <eos>).

- Attention consistently focuses on 'miss' and 'me?' across all three runs, confirming stable alignment learning.
- Steps generating the pronoun 'je' attend to earlier tokens; steps generating 'manqué' attend strongly to 'miss' and 'me?'.
- The consistent red/dark pattern at the 'miss' and 'me?' columns confirms the model correctly identifies the most semantically relevant source tokens at each decode step.

This stable alignment pattern across multiple runs is strong evidence that the Bahdanau attention mechanism has learned meaningful EN-FR word-level correspondences purely from parallel data, without any explicit alignment supervision.